



Cisco OpenStack Installer Deployment Guide, Grizzly Release

First Published: June 25, 2013

Americas Headquarters

Cisco Systems, Inc.
170 West Tasman Drive
San Jose, CA 95134-1706
USA
<http://www.cisco.com>
Tel: 408 526-4000
800 553-NETS (6387)
Fax: 408 527-0883

THE SPECIFICATIONS AND INFORMATION REGARDING THE PRODUCTS IN THIS MANUAL ARE SUBJECT TO CHANGE WITHOUT NOTICE. ALL STATEMENTS, INFORMATION, AND RECOMMENDATIONS IN THIS MANUAL ARE BELIEVED TO BE ACCURATE BUT ARE PRESENTED WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED. USERS MUST TAKE FULL RESPONSIBILITY FOR THEIR APPLICATION OF ANY PRODUCTS.

THE SOFTWARE LICENSE AND LIMITED WARRANTY FOR THE ACCOMPANYING PRODUCT ARE SET FORTH IN THE INFORMATION PACKET THAT SHIPPED WITH THE PRODUCT AND ARE INCORPORATED HEREIN BY THIS REFERENCE. IF YOU ARE UNABLE TO LOCATE THE SOFTWARE LICENSE OR LIMITED WARRANTY, CONTACT YOUR CISCO REPRESENTATIVE FOR A COPY.

The Cisco implementation of TCP header compression is an adaptation of a program developed by the University of California, Berkeley (UCB) as part of UCB's public domain version of the UNIX operating system. All rights reserved. Copyright © 1981, Regents of the University of California.

NOTWITHSTANDING ANY OTHER WARRANTY HEREIN, ALL DOCUMENT FILES AND SOFTWARE OF THESE SUPPLIERS ARE PROVIDED "AS IS" WITH ALL FAULTS. CISCO AND THE ABOVE-NAMED SUPPLIERS DISCLAIM ALL WARRANTIES, EXPRESSED OR IMPLIED, INCLUDING, WITHOUT LIMITATION, THOSE OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OR ARISING FROM A COURSE OF DEALING, USAGE, OR TRADE PRACTICE.

IN NO EVENT SHALL CISCO OR ITS SUPPLIERS BE LIABLE FOR ANY INDIRECT, SPECIAL, CONSEQUENTIAL, OR INCIDENTAL DAMAGES, INCLUDING, WITHOUT LIMITATION, LOST PROFITS OR LOSS OR DAMAGE TO DATA ARISING OUT OF THE USE OR INABILITY TO USE THIS MANUAL, EVEN IF CISCO OR ITS SUPPLIERS HAVE BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

Any Internet Protocol (IP) addresses and phone numbers used in this document are not intended to be actual addresses and phone numbers. Any examples, command display output, network topology diagrams, and other figures included in the document are shown for illustrative purposes only. Any use of actual IP addresses or phone numbers in illustrative content is unintentional and coincidental.

Cisco and the Cisco logo are trademarks or registered trademarks of Cisco and/or its affiliates in the U.S. and other countries. To view a list of Cisco trademarks, go to this URL: <http://www.cisco.com/go/trademarks>. Third-party trademarks mentioned are the property of their respective owners. The use of the word partner does not imply a partnership relationship between Cisco and any other company. (1110R)

© 2013 Cisco Systems, Inc. All rights reserved.



CONTENTS

Preface

Preface v

Audience v

Document Conventions v

Related Documentation vii

Obtaining Documentation and Submitting a Service Request vii

CHAPTER 1

Overview 1

About Cisco OpenStack Installer 1

Supported Deployments 1

Guidelines and Limitations of Cisco OSI 2

CHAPTER 2

System Prerequisites 5

System Requirements for Cisco OSI 5

Supported Hypervisors 7

Information Needed to Deploy Cisco OSI 7

Sample site.pp.example File 14

CHAPTER 3

Installing OpenStack 21

Summary of Steps for Installing OpenStack with Cisco OSI 21

Cisco OSI Puppet Manifest Files 22

Creating the Build Server 23

Creating the Build Server with the Cisco OSI Install Script 23

Creating the Build Server Manually 24

Customizing the Build Server 26

Building the Control and Compute Nodes 27

Building the Control and Compute Nodes Individually 27

Building the Control and Compute Nodes with the clean_node.sh Script 28

Building the Control and Compute Nodes with the reset_nodes.sh Script	28
Rerunning Puppet on the Build Node and Enabling Puppet to Run as an Agent	29
Testing the OpenStack Deployment	29

CHAPTER 4**Using OpenStack 31**

Deploying the First Virtual Machine Through a Script	31
Monitoring the Health of your System in Nagios	32
Viewing Statistics in Graphite	32
Additional Documentation	33



Preface

This preface contains the following sections:

- [Audience, page v](#)
- [Document Conventions, page v](#)
- [Related Documentation, page vii](#)
- [Obtaining Documentation and Submitting a Service Request, page vii](#)

Audience

This guide is intended primarily for experienced data center and/or network administrators who want to use Cisco OpenStack Installer (Cisco OSI) to deploy an OpenStack cluster.

Document Conventions

This document uses the following conventions:

Convention	Description
^ or Ctrl	Both the ^ symbol and Ctrl represent the Control (Ctrl) key on a keyboard. For example, the key combination ^D or Ctrl-D means that you hold down the Control key while you press the D key. (Keys are indicated in capital letters but are not case sensitive.)
bold font	Commands and keywords and user-entered text appear in bold font .
<i>Italic font</i>	Document titles, new or emphasized terms, and arguments for which you supply values are in <i>italic font</i> .
Courier font	Terminal sessions and information the system displays appear in <code>courier font</code> .
Bold Courier font	Bold Courier font indicates text that the user must enter.

Convention	Description
[x]	Elements in square brackets are optional.
...	An ellipsis (three consecutive nonbolded periods without spaces) after a syntax element indicates that the element can be repeated.
	A vertical line, called a pipe, indicates a choice within a set of keywords or arguments.
[x y]	Optional alternative keywords are grouped in brackets and separated by vertical bars.
{x y}	Required alternative keywords are grouped in braces and separated by vertical bars.
[x {y z}]	Nested set of square brackets or braces indicate optional or required choices within optional or required elements. Braces and a vertical bar within square brackets indicate a required choice within an optional element.
string	A nonquoted set of characters. Do not use quotation marks around the string or the string will include the quotation marks.
< >	Nonprinting characters such as passwords are in angle brackets.
[]	Default responses to system prompts are in square brackets.
!, #	An exclamation point (!) or a pound sign (#) at the beginning of a line of code indicates a comment line.

Reader Alert Conventions

This document uses the following conventions for reader alerts:



Note

Means *reader take note*. Notes contain helpful suggestions or references to material not covered in the manual.



Tip

Means *the following information will help you solve a problem*.



Caution

Means *reader be careful*. In this situation, you might do something that could result in equipment damage or loss of data.

**Timesaver**

Means *the described action saves time*. You can save time by performing the action described in the paragraph.

**Warning**

Means *reader be warned*. In this situation, you might perform an action that could result in bodily injury.

Related Documentation

The following documentation provides additional information about customizing, configuring, and maintaining your OpenStack cluster:

- [Cisco OpenStack documentation wiki](#)
- [OpenStack documentation](#) available on the OpenStack website

Obtaining Documentation and Submitting a Service Request

For information on obtaining documentation, submitting a service request, and gathering additional information, see the monthly *What's New in Cisco Product Documentation*, which also lists all new and revised Cisco technical documentation, at:

<http://www.cisco.com/en/US/docs/general/whatsnew/whatsnew.html>

Subscribe to the *What's New in Cisco Product Documentation* as a Really Simple Syndication (RSS) feed and set content to be delivered directly to your desktop using a reader application. The RSS feeds are a free service and Cisco currently supports RSS version 2.0.



Overview

This chapter contains the following sections:

- [About Cisco OpenStack Installer, page 1](#)
- [Supported Deployments, page 1](#)
- [Guidelines and Limitations of Cisco OSI, page 2](#)

About Cisco OpenStack Installer

Cisco OpenStack Installer (Cisco OSI) simplifies the OpenStack installation process through a series of Puppet modules with a small number of Bash scripts. Cisco OSI enables you to deploy a multinode OpenStack system with Quantum-enabled network services management and a simple monitoring function based on Nagios, collectd, and Graphite.

When you deploy OpenStack with Cisco OSI, you configure an initial build server outside the OpenStack cluster to manage and automate the OpenStack software deployment. During the installation, the build server primarily functions in the following ways:

- As a Puppet Master server that deploys the software onto and manages the configuration of the OpenStack cluster. For more information about Puppet, see the [Puppet Labs documentation](#).
- As a Cobbler installation server to manage the Pre-boot Execution Environment (PXE) boot that is used for rapid bootstrapping of the OpenStack cluster. For more information about Cobbler, see the [Cobbler user documentation](#).

After the build server is installed and configured, it is used as an out-of-band automation and management workstation to bring up, control, and reconfigure (if necessary) the nodes of the OpenStack cluster. It also functions as a monitoring server to collect statistics about the health and performance of the OpenStack cluster, as well as to monitor the availability of the servers and services of the OpenStack cluster.

Supported Deployments

The instructions in this installation guide deploy the following environment with Cisco OSI.

- Single build node

- Single control node
- Single network node
- Multiple compute nodes
- Quantum-managed network in an OpenVirtual Switch (OVS) Generic Routing Encapsulation (GRE) tunnel-based environment

Cisco OSI deploys several OpenStack components and some non-OpenStack components, including the following :

- Nova
- Keystone
- Cinder
- Horizon
- Quantum
- Swift
- Nagios
- collectd
- Graphite
- Other OpenStack community incubated projects

Cisco OSI does not deploy any other OpenStack components, such as the following components:

- Ceilometer
- Other OpenStack community incubated projects

Guidelines and Limitations of Cisco OSI

Cisco OSI installation and configuration guidelines and limitations are as follows:

IP Networks

The following IP networks are used by Cisco OSI:

- Management IP network
- Floating IP network
- Instance IP network (vm_net) used to pass private tenant network traffic through GRE tunnels

Proxy Configurations

If your network requires that you use a proxy server to access the internet, you might encounter issues in the phase of the installation process when Cisco OSI is downloading and installing software packages. Proxy server configuration issues can result in the following problems:

- **apt** reports hash mismatches or file corruptions when it verifies the downloaded files.
- **apt-get** reports a hash mismatch when an issue occurs with a caching engine. If possible, you should bypass the caching engine to resolve this problem.
- PIP fails to install packages correctly. The PIP component installs itself locally. If you have a manually defined proxy on your build node, the PIP component will attempt to go through the proxy server to get back to its local resources and will fail.

If your environment includes a proxy server, we recommend that you do the following:

- Export the two types of proxies needed in your root shell when you run **fetch** commands, as noted in the relevant procedures.
- Configure the **\$proxy** setting in the `site.pp` to reflect your local proxy server.
- If all the servers that host the build, control, and compute nodes in your OpenStack cluster do not have public Internet accessible IPs and you are building the cluster in a controlled environment, ensure that the **\$default_gateway** setting is not configured in the `site.pp` file and all the files that are required to install the control and compute nodes are fetched from the boot server.

In addition, if you have proxies, and you set your proxy information in either your `.profile` or in a file such as `/etc/environment`, you must set both `http_proxy` and `https_proxy`. You will also need to set a **no_proxy** command at least for the build node. The required settings for proxy information are as follows:

```
http_proxy=http://your-proxy.address.com:80/  
https_proxy=https://your-https-proxy.address.com:443/  
no_proxy=your-build-node-name,*yourbuild.domain.name,127.0.0.1,127.0.1.1,localhost
```

Supported Servers

Cisco OSI has been tested against a system that includes Cisco UCS servers. Systems with non-supported servers may encounter issues.



System Prerequisites

This chapter contains the following sections:

- [System Requirements for Cisco OSI, page 5](#)
- [Supported Hypervisors, page 7](#)
- [Information Needed to Deploy Cisco OSI, page 7](#)

System Requirements for Cisco OSI

Supported Servers

Cisco OSI supports deployment of OpenStack on Cisco Unified Computing System (Cisco UCS) servers. Several steps in the automation provided by Cisco OSI leverage the Cisco UCS management plane (either Cisco UCS Manager or Cisco Integrated Management Controller) to execute system tasks.

You can deploy Cisco OSI on the following Cisco UCS servers:

- Standalone rack-mount servers that are managed by Cisco Integrated Management Controller
- Integrated rack-mount servers that are managed by Cisco UCS Manager
- Blade servers that are managed by Cisco UCS Manager

Recommended Release Levels

The following release levels are recommended for any server that is part of an OpenStack cluster:

- For blade servers and integrated rack-mount servers, Cisco UCS Manager, Release 2.1(1) and later
- For standalone rack-mount servers, Cisco Integrated Management Controller, Release 1.5 and later

Minimum Server Requirements

The following table lists the minimum requirements for the Cisco UCS servers that you use for the nodes in your OpenStack cluster:

**Tip**

Additional information about the minimum requirements for the control and compute nodes is available in the [Compute and Image System Requirements section](#) of the *OpenStack Compute Administration Guide* for the Grizzly release.

Server/Node	Recommended Hardware	Notes
Build node	Memory: 4 GB (RAM) Disk space: 20 GB	The build node must also have Internet connectivity to be able to download Cisco OSI modules and manifests. To ensure that the build node can build and communicate with the other nodes in your cluster, it must also have a network interface on the same network as the management interfaces of the other OpenStack cluster servers.
Control node	Processor: 64-bit x86 Memory: 12 GB (RAM) Disk space: 1 TB (SATA or SAS) Network: One 1 Gbps Network Interface Card (NIC)	Two NICs are recommended but not required. A quad core server with 12 GB of RAM is sufficient for a control node.
Compute node	Processor: 64-bit x86 Memory: 128 GB (RAM) Disk space: 300 GB (SATA) Volume storage: two disks with 2 TB (SATA) for volumes attached to the compute nodes Network: Two 1 Gbps NICs	
Proxy node	Processor: 64-bit x86 Memory: 12 GB (RAM) Disk space: 1 TB (SATA or SAS) Network: One 1 Gbps Network Interface Card (NIC)	

Server/Node	Recommended Hardware	Notes
Storage node	Processor: 64-bit x86 Memory: 32 GB (RAM) Disk space: 300 GB (SATA) Volume storage: <ul style="list-style-type: none"> For rack-mount servers, either 24 disks with 1 TB (SATA) or 12 disks with 3TB (SATA) depending upon the model For blade servers, 2 disks with 1 TB (SATA) for combined base OS and storage Network: Two 1 Gbps NICs	Three or more storage nodes are needed.

Supported Operating Systems

Cisco OSI supports and deploys the Ubuntu 12.04 LTS operating system.

Supported Hypervisors

Cisco OSI supports the Kernel-based Virtual Machine (KVM) for the supported Ubuntu operating system as the hypervisor in the OpenStack cluster. KVM is a virtualization package for Linux on an x86 hardware platform. KVM uses x86 hardware virtualization extensions (for example, Intel VT-x) to implement a hypervisor that hosts VMs as userspace processes.

Information Needed to Deploy Cisco OSI

When you deploy Cisco OSI, you must provide the following information when you configure the `site.pp` file to customize the build server for your OpenStack environment. You can use the default values for all other settings included in the `site.pp` file.

The `site.pp` file contains the configuration settings for and information about the nodes in your OpenStack deployment. Cisco OSI uses these settings to configure the nodes during installation and deployment:

Proxy Configuration

If your network includes an HTTP/HTTPS proxy, you must define the proxy URL. Modify the `$proxy` variable with the following information:

- Proxy server
- Proxy port number

Package Repository Configuration

The package repository used to install OpenStack. If necessary, modify one of the following variables:

- *\$package_repo*—By default, the value of this variable is `cisco_repo`. If you prefer to use the upstream Ubuntu package from the cloud archive, use `cloud_archive` for the value.
- *\$location*—By default, this points to the Cisco FTP distribution:
`ftp://ftpeng.cisco.com/openstack/cisco`. If your system includes a proxy server, you can use the Cisco HTTP distribution location:
`http://openstack-repo.cisco.com/openstack/cisco`. However, the HTTP location is not permanent and might change at any time.

Build Node Name

Modify the *\$build_node_name* variable with the hostname of the build server.

All letters in the build node name must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.

NTP Server Name

Modify the *\$ntp_servers* variable with the fully qualified domain name/location of a time server that is accessible to the build server. An accessible DNS server must be able to resolve the server name that you enter here.

Build Node Cobbler Variables

The network settings for the build server that are required for the Cobbler installation server that manages the PXE boot. You must define the following variables:

- *\$cobbler_node_ip*—IP address (IPv4) of the build node.
- *\$node_subnet*—Subnet on which the build node will be located.
- *\$node_netmask*—Netmask of the build node.
- *\$node_gateway*—Gateway. If the network does not include a default route through a gateway, comment out this line. If you do so, the files that are required for installing the control and compute nodes are fetched from the boot server.
- *\$domain_name*—Domain name. The build node and compute nodes use this domain name for the local DNS that serves addresses in the domain.

Preseed File Configuration

The settings used in a preseed file called `cisco-preseed` in `/etc/cobbler/preseeds/`. The preseed file is used to install Ubuntu on the OpenStack nodes. We recommend that you use the default settings. However, if HTTPS is disabled and Cisco UCS Manager uses HTTP, you must change the `ucsm_port` setting.

These settings include the following variables:

- `$admin_user`—The default admin username is `localadmin`.
- `$password_crypted`—This is an SHA-512 hashed password. See the notes in the `site.pp` file for the default password and instructions on how to generate a new password.
- `$autostart_puppet`—Whether the Puppet agent will start automatically.
- `$ucsm_port`—For OpenStack clusters that use servers managed by Cisco UCS Manager. By default, this setting configures communication on HTTPS, using port 443.

OpenStack Variables

The settings used to deploy and configure OpenStack after Ubuntu has been installed on the OpenStack nodes. You must define the following variables:

- `$controller_node_address`—IP address of the control node.
- `$controller_node_network`—Network address for the control node.
- `$controller_hostname`—Controller hostname. All letters in the control node hostname must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.
- `$db_allowed_network`—Address for a network that has access to the MySQL database on the control node. This address should be the same network as the one for the control node, but this network address should include the MySQL network wildcard syntax. For example, `209.165.202.%`
- `$swift_proxy_address`—Address of the Swift proxy node. If you have multiple Swift proxy nodes, this address should be the virtual IP address (VIP) used to load-balance across the individual nodes.
- `$public_interface`—Public interface of the node. The default value for this setting is `eth0`. This interface must have an IP address reachable by all other nodes in the OpenStack cluster. This address is used for API access, for the Horizon user interface, and as an endpoint for the default GRE tunnel mechanism used in the OVS network configuration.
- `$external_interface`—External interface of the node. The default value for this setting is `eth1`. This interface is only required on the network/control node. It provides a Layer 2 path for the `l3_agent` external router interface. This interface must be attached to an upstream device that provides a Layer 3 router interface, with the default router configuration assuming that the first non "network" address in the external network IP subnet is the default forwarding path if no more specific host routes are added.
- `$install_drive`—Drive on which Ubuntu and OpenStack are installed on each node. The default value for this setting is `/dev/sda`. The assumption is that all nodes are installed on the same device name.

OpenStack Service Credentials

The settings used to change the usernames and passwords for the OpenStack services. We recommend that you change the following defaults for all production deployments:

- *\$admin_email*—Admin email
- *\$admin_password*—Admin password
- *\$keystone_db_password*—Keystone database password
- *\$keystone_admin_token*—Keystone admin token
- *\$nova_user*—Nova username
- *\$nova_db_password*—Nova database password
- *\$nova_user_password*—Nova user password
- *\$libvirt_type*—Type of libvirt used for the OpenStack cluster. You do not need to change the default value of kvm
- *\$glance_db_password*—Glance database
- *\$glance_user_password*—Glance user password
- *\$glance_sql_connection*—Glance SQL connection
- *\$glance_on_swift*—Glance on Swift. This value can be true or false
- *\$cinder_user_password*—Cinder user password
- *\$cinder_db_password*—Cinder database password
- *\$quantum_user_password*—Quantum user password
- *\$quantum_db_password*—Quantum database password
- *\$rabbit_password*—Rabbit password
- *\$rabbit_user*—Rabbit username
- *\$swift_password*—Swift password
- *\$swift_user*—Swift username
- *\$swift_hash*—Swift hash value
- *\$sql_connection*—Nova database connection
- *\$glance_backend*—Glance backend configuration. This supports file or swift

Storage Configuration

The definitions required to configuration storage for Cinder. You must review and define the following variables:

- *\$cinder_controller_enabled*—The default value is true, which enables Cinder services.
- *\$cinder_compute_enabled*—The default value is true, which deploys Cinder on all compute nodes.
- *\$cinder_storage_driver*—The Cinder storage driver used by the OpenStack cluster. The default value is iscsi. See the `site.pp` file for other storage drivers, such as netapp or nfs.
-

OpenStack Control Node Definitions

The definitions required for the control node. The `site.pp.example` file contains sample definitions for different control node configurations. You must define the following variables:

- *cobbler_node*—Control node hostname. All letters in the control node hostname must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.
- *node_type*—Node type. This value must be control.
- *mac*—MAC address for the boot interface of the control node.
- *ip*—IP address of the control node.
- *power_address*—Power address. The IP address of Cisco UCS Manager or Cisco Integrated Management Controller that manages the server that acts as the control node. For Cisco UCS Manager in a high availability or cluster configuration, the power address is the cluster IP address.
- *power_id*—Optional. The service profile associated with the Cisco UCS server. This value is only required if the server has a *power_address* that is configured with a subgroup in Cisco UCS Manager.
- *power_type*—By default, this variable is not specified and assumes the value of ipmitool that is required for standalone rack-mount servers. Therefore, you do not must specify this variable if your cluster uses standalone rack-mount servers. However, if your cluster uses blade servers and/or rack-mount servers managed by Cisco UCS Manager, this must have a value of UCS.
- *power_user*—The username required to log in for power management of the server.
- *power_password*—The password required to log in for power management of the server.

OpenStack Compute Node Definitions

The definitions required for the compute node. The `site.pp` file includes a block for your first compute node. If your OpenStack deployment includes more than one compute node, copy this block, add an additional block to the `site.pp` file for each compute node, and update the definitions as needed.

The `site.pp.example` file contains sample definitions for different compute node configurations. The compute node variables are as follows:

- *cobbler_node*—Compute node hostname. All letters in the compute node hostname must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.
- *node_type*—Node type. This value must be `compute`.
- *mac*—MAC address for the boot interface of the compute node.
- *ip*—IP address of the server that acts as this node.
- *power_address*—Power address. The IP address of Cisco UCS Manager or Cisco Integrated Management Controller that manages the server that acts as the compute node. For Cisco UCS Manager in a high availability or cluster configuration, the power address is the cluster IP address.
- *power_id*—Optional. The service profile associated with the Cisco UCS server. This value is only required if the server has a *power_address* that is configured with a subgroup in Cisco UCS Manager.
- *power_type*—By default, this variable is not specified and assumes the value of `ipmitool` that is required for standalone rack-mount servers. Therefore, you do not must specify this variable if your cluster uses standalone rack-mount servers. However, if your cluster uses blade servers and/or rack-mount servers managed by Cisco UCS Manager, this must have a value of `UCS`.
- *power_user*—The username required to log in for power management of the server.
- *power_password*—The password required to log in for power management of the server.

Swift Proxy Node Cobbler Definitions

The definitions required by Cobbler for the first Swift proxy node. The variables for this node are as follows:

- *cobbler_node*—Swift proxy node hostname. All letters in the hostname must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.
- *node_type*—Node type. This value must be `swift-proxy`.
- *mac*—MAC address for the boot interface of the Swift proxy node.
- *ip*—IP address of the server that acts as this node.
- *power_address*—Power address. The IP address of Cisco UCS Manager or Cisco Integrated Management Controller that manages the server that acts as the Swift proxy node. For Cisco UCS Manager in a high availability or cluster configuration, the power address is the cluster IP address.

Swift Storage Node Cobbler Definitions

The definitions required by Cobbler for the first Swift storage node. If your OpenStack deployment includes more than one storage node, copy this block, add an additional block to the `site.pp` file for each storage node, and update the definitions as needed

The variables for this node are as follows:

- *cobbler_node*—Swift storage node hostname. All letters in the hostname must be in lowercase. For more information about this Puppet limitation, see <http://projects.puppetlabs.com/issues/1168.html>.
- *node_type*—Node type. This value must be `swift-storage`.
- *mac*—MAC address for the boot interface of the Swift storage node.
- *ip*—IP address of the server that acts as this node.
- *power_address*—Power address. The IP address of Cisco UCS Manager or Cisco Integrated Management Controller that manages the server that acts as the Swift storage node. For Cisco UCS Manager in a high availability or cluster configuration, the power address is the cluster IP address.

Node Types

The settings that specify the hostnames of the servers for each node in your OpenStack deployment. If your OpenStack deployment includes more than one compute node, copy the compute node block and add an additional block to the `site.pp` file for each compute node.

These hostnames are the same hostnames that you specified for each of the nodes, and must all be in lowercase due to the Puppet limitation.

Swift Proxy Node Puppet Definitions

The definitions required by Puppet for the Swift proxy node. The variables for these definitions are as follows:

- *swift_local_net_ip*—IP address of the server that acts as the first Swift proxy node.
- *keystone_host*—Address of the control node. This value of `$controller_node_address` defaults to the address of the control node.
- *swift_user_password*—User password for Swift.
- *swift_hash_suffix*—Hash suffix used by Swift.

Swift Storage Node Puppet Definitions

The definitions required by Puppet for the first Swift storage node. If your OpenStack deployment includes more than one storage node, copy this block, add an additional block to the `site.pp` file for each storage node, and update the definitions as needed

The variables for this node are as follows:

- `swift_zone`—The zone to which this storage node belongs.
- `swift_local_net_ip`—IP address of the server that acts as this node.
- `storage_type`—Type of storage used by this node. The default value is disk.
- `storage_devices`—The list of disk devices to be formatted with XFS and used for Swift storage.
- `swift_hash_suffix`—Hash suffix used by Swift.

Sample site.pp.example File

```
# This document serves as an example of how to deploy
# basic multi-node openstack environments.
# In this scenario Quantum is using OVS with GRE Tunnels

##### Proxy Configuration #####
# If you use an HTTP/HTTPS proxy, uncomment this setting and specify the
# correct proxy URL. If you do not use an HTTP/HTTPS proxy, leave this
# setting commented out.
#$proxy = "http://proxy-server:port-number"

##### package repo configuration #####
#
# The package repos used to install openstack
$package_repo = 'cisco_repo'
# Alternatively, the upstream Ubuntu package from cloud archive can be used
# $package_repo = 'cloud_archive'

# If you are behind a proxy you may choose not to use our ftp distribution, and
# instead try our http distribution location. Note the http location is not
# a permanent location and may change at any time.
$location = "ftp://ftpeng.cisco.com/openstack/cisco"
# Alternate, uncomment this one, and comment out the one above.
#$location = "http://openstack-repo.cisco.com/openstack/cisco"

##### Build Node (Cobbler, Puppet Master, NTP) #####
# Change the following to the host name you have given your build node.
# This name should be in all lower case letters due to a Puppet limitation
# (refer to http://projects.puppetlabs.com/issues/1168).
$build_node_name = "build-server"

##### NTP Configuration #####
# Change this to the location of a time server in your organization accessible
# to the build server. The build server will synchronize with this time
# server, and will in turn function as the time server for your OpenStack
# nodes.
$ntp_servers = ["time-server.domain.name"]

##### Build Node Cobbler Variables #####
# Change these 5 parameters to define the IP address and other network
# settings of your build node. The cobbler node *must* have this IP
# configured and it *must* be on the same network as the hosts to install.
$cobbler_node_ip = '192.168.242.100'
$node_subnet = '192.168.242.0'
$node_netmask = '255.255.255.0'
```

```
# This gateway is optional - if there's a gateway providing a default route,
# specify it here. If not, comment this line out.
$node_gateway = '192.168.242.1'
# This domain name will be the name your build and compute nodes use for the
# local DNS. It doesn't have to be the name of your corporate DNS - a local
# DNS server on the build node will serve addresses in this domain - but if
# it is, you can also add entries for the nodes in your corporate DNS
# environment they will be usable *if* the above addresses are routeable
# from elsewhere in your network.
$domain_name = 'domain.name'
# This setting likely does not need to be changed.
# To speed installation of your OpenStack nodes, it configures your build
# node to function as a caching proxy storing the Ubuntu install files used
# to deploy the OpenStack nodes.
$cobbler_proxy = "http://${cobbler_node_ip}:3142/"

##### Preseed File Configuration #####
# This will build a preseed file called 'cisco-preseed' in
# /etc/cobbler/preseeds/
# The preseed file automates the installation of Ubuntu onto the OpenStack
# nodes.
#
# The following variables may be changed by the system admin:
# 1) admin_user
# 2) password_crypted
# 3) autostart_puppet -- whether the puppet agent will auto start
# Default user is: localadmin
# Default SHA-512 hashed password is "ubuntu":
#
#$UfgWxrIv$k4KfzAEMqMg.fppmSOTd0usI4j6gfjs0962.JXsoJRwa5wMz8yQk4SfInn4.WZ3L/MCt5u.62tHDGB36EhiKF1
# To generate a new SHA-512 hashed password, run the following replacing
# the word "password" with your new password. Then use the result as the
# $password_crypted variable.
# python -c "import crypt, getpass, pwd; print crypt.crypt('password', '\$6\$UfgWxrIv\$')"
$admin_user = 'localadmin'
$password_crypted =
'$6$UfgWxrIv$k4KfzAEMqMg.fppmSOTd0usI4j6gfjs0962.JXsoJRwa5wMz8yQk4SfInn4.WZ3L/MCt5u.62tHDGB36EhiKF1'
$autostart_puppet = true

# If the setup uses the UCS B-series blades, enter the port on which the
# ucsd accepts requests. By default the UCSM is enabled to accept requests
# on port 443 (https). If https is disabled and only http is used, set
# $ucsm_port = '80'
$ucsm_port = '443'

##### OpenStack Variables #####
# These values define parameters which will be used to deploy and configure
# OpenStack once Ubuntu is installed on your nodes.
#
# Change these next 3 parameters to the network settings of the node which
# will be your OpenStack control node. Note that the $controller_hostname
# should be in all lowercase letters due to a limitation of Puppet
# (refer to http://projects.puppetlabs.com/issues/1168).
$controller_node_address = '192.168.242.10'
$controller_node_network = '192.168.242.0'
$controller_hostname = 'control-server'
# Specify the network which should have access to the MySQL database on
# the OpenStack control node. Typically, this will be the same network as
# defined in the controller node_network parameter above. Use MySQL network
# wild card syntax to specify the desired network.
$db_allowed_network = '192.168.242.%'
# These next two values typically do not need to be changed. They define the
# network connectivity of the OpenStack controller. This is the interface
# used to connect to Horizon dashboard.
$controller_node_public = $controller_node_address
# This is the interface used for external backend communication.
$controller_node_internal = $controller_node_address

# Specify the address of the Swift proxy
# If you have multiple Swift proxy nodes, this should be the address
# of the VIP used to load-balance across the individual nodes
$swift_proxy_address = '192.168.242.179'
```

```

# These next two parameters specify the networking hardware used in each node
# Current assumption is that all nodes have the same network interfaces and are
# cabled identically. However, with the control node acting as network node,
# only the control node requires 2 interfaces. For all other nodes, a single
# interface is functional with the assumption that:
#   a) The public_interface will have an IP address reachable by
#       all other nodes in the openstack cluster. This address will
#       be used for API Access, for the Horizon UI, and as an endpoint
#       for the default GRE tunnel mechanism used in the OVS network
#       configuration.
#   b) The external_interface is used to provide a Layer2 path for
#       the l3_agent external router interface. It is expected that
#       this interface be attached to an upstream device that provides
#       a L3 router interface, with the default router configuration
#       assuming that the first non "network" address in the external
#       network IP subnet will be used as the default forwarding path
#       if no more specific host routes are added.
#
# It is assumed that this interface has an IP Address associated with it
# and is available and connected on every host in the OpenStack cluster
$public_interface = 'eth0'
# The external_interface is used for external connectivity in association
# with the l3_agent external router interface, providing floating IPs
# (this is only required on the network/controller node)
$external_interface = 'eth1'

# Select the drive on which Ubuntu and OpenStack will be installed in each
# node. The current assumption is that all nodes will be installed on the
# same device name.
$install_drive = '/dev/sda'

##### OpenStack Service Credentials #####
# This block of parameters is used to change the user names and passwords
# used by the services which make up OpenStack. The following defaults should
# be changed for any production deployment.
$admin_email = 'root@localhost'
$admin_password = 'Cisco123'
$keystone_db_password = 'keystone_db_pass'
$keystone_admin_token = 'keystone_admin_token'
$nova_user = 'nova'
$nova_db_password = 'nova_pass'
$nova_user_password = 'nova_pass'
$libvirt_type = 'kvm'
$glance_db_password = 'glance_pass'
$glance_user_password = 'glance_pass'
$glance_sql_connection =
"mysql://glance:${glance_db_password}@${controller_node_address}/glance"
$glance_on_swift = false
$cinder_user_password = 'cinder_pass'
$cinder_db_password = 'cinder_pass'
$quantum_user_password = 'quantum_pass'
$quantum_db_password = 'quantum_pass'
$rabbit_password = 'openstack_rabbit_password'
$rabbit_user = 'openstack_rabbit_user'
$swift_password = 'openstack_swift_password'
$swift_user = 'openstack_swift_user'
$swift_hash = 'swift_secret'
# Nova DB connection
$sql_connection = "mysql://${nova_user}:${nova_db_password}@${controller_node_address}/nova"
# Glance backend configuration, supports 'file' or 'swift'.
$glance_backend = 'file'

##### Test variables #####
# Variables used to populate test script:
# /tmp/test_nova.sh
#
# Image to use for tests. Accepts 'kvm' or 'cirros'.
$test_file_image_type = 'kvm'

#### end shared variables #####

# Storage Configuration
# Set to true to enable Cinder services.

```



```

$cinder_controller_enabled      = true

# Set to true to enable Cinder deployment to all compute nodes.
$cinder_compute_enabled        = true

# The cinder storage driver to use. Default is 'iscsi'.
$cinder_storage_driver          = 'iscsi'

# Other drivers exist for cinder. Here are examples on how to enable them.
#
# NetApp iSCSI Driver
# $cinder_storage_driver = 'netapp'
# $netapp_wsdl_url        = ''
# $netapp_login           = ''
# $netapp_password        = ''
#
# NFS
# share information is stored in flat text file specified in
# $nfs_shares_config. The format for this file is "hostname:/mountpoint",
# e.g. "192.168.2.55:/myshare", with only one entry per line.
#
# $cinder_storage_driver = 'nfs'
# $nfs_shares_config      = '/etc/cinder/shares.conf'

##### OpenStack Node Definitions #####
# This section is used to define the hardware parameters of the nodes
# which will be used for OpenStack. Cobbler will automate the installation
# of Ubuntu onto these nodes using these settings.

# The build node name is changed in the "node type" section further down
# in the file. This line should not be changed here.
node 'build-node' inherits master-node {

# This block defines the control server. Replace "control-server" with the
# host name of your OpenStack controller, and change the "mac" to the MAC
# address of the boot interface of your OpenStack controller. Change the
# "ip" to the IP address of your OpenStack controller. The power_address
# parameter specifies the address to use for device power management,
# power_user and power_password specify the login credentials to use for
# power management, and power_type determines which Cobbler fence script
# is used for power management. Supported values for power_type are
# 'ipmitool' for generic IPMI devices and UCS C-series servers in standalone
# mode or 'ucs' for C-series or B-series UCS devices managed by UCSM.

  cobbler_node { "control-server":
    node_type => "control",
    mac => "00:11:22:33:44:55",
    ip => "192.168.242.10",
    power_address => "192.168.242.110",
    power_user => "admin",
    power_password => "password",
    power_type => "ipmitool"
  }

# This block defines the first compute server. Replace "compute-server01"
# with the host name of your first OpenStack compute node (note: the hostname
# should be in all lowercase letters due to a limitation of Puppet; refer to
# http://projects.puppetlabs.com/issues/1168), and change the "mac" to the
# MAC address of the boot interface of your first OpenStack compute node.
# Change the "ip" to the IP address of your first OpenStack compute node.

# Begin compute node
  cobbler_node { "compute-server01":
    node_type => "compute",
    mac => "11:22:33:44:55:66",
    ip => "192.168.242.21",
    power_address => "192.168.242.121",
    power_user => "admin",
    power_password => "password",
    power_type => "ipmitool"
  }
}

```

```

# Example with UCS blade power_address with a sub-group (in UCSM), and
# a ServiceProfile for power_id.
# cobbler_node { "compute-server01":
#   node_type => "compute",
#   mac => "11:22:33:44:66:77",
#   ip => "192.168.242.21",
#   power_address => "192.168.242.121:org-cisco",
#   power_id => "OpenStack-1",
#   power_type => 'ucs',
#   power_user => 'admin',
#   power_password => 'password'
# }
# End compute node

### Repeat as needed ###
# Make a copy of your compute node block above for each additional OpenStack
# node in your cluster and paste the copy in this section. Be sure to change
# the host name, mac, ip, and power settings for each node.

# This block defines the first swift proxy server. Replace "swift-server01"
# with the host name of your first OpenStack swift proxy node (note:
# the hostname should be in all lowercase letters due to a limitation of
# Puppet; refer to http://projects.puppetlabs.com/issues/1168), and change
# the "mac" to the MAC address of the boot interface of your first OpenStack
# swift proxy node. Change the "ip" to the IP address of your first
# OpenStack swift proxy node.

# Begin swift proxy node
# cobbler_node { "swift-proxy01":
#   node_type => "swift-proxy",
#   mac => "11:22:33:aa:bb:cc",
#   ip => "192.168.242.179",
#   power_address => "192.168.242.12"
# }

# This block defines the first swift storage server. Replace "swift-storage01"
# with the host name of your first OpenStack swift storage node (note: the
# hostname should be in all lowercase letters due to a limitation of Puppet;
# refer to http://projects.puppetlabs.com/issues/1168), and change the "mac"
# to the MAC address of the boot interface of your first OpenStack swift
# storage node. Change the "ip" to the IP address of your first OpenStack
# swift storage node.

# Begin swift storage node
# cobbler_node { "swift-storage01":
#   node_type => "swift-storage",
#   mac => "11:22:33:cc:bb:aa",
#   ip => "192.168.242.180",
#   power_address => "192.168.242.13"
# }

### Repeat as needed ###
# Make a copy of your swift storage node block above for each additional
# node in your swift cluster and paste the copy in this section. Be sure
# to change the host name, mac, ip, and power settings for each node.

### End repeated nodes ###
}

### Node types ###
# These lines specify the host names in your OpenStack cluster and what the
# function of each host is. internal_ip should be the same as what is
# specified as "ip" in the OpenStack node definitions above.
# This sets the IP for the private(internal) interface of controller nodes
# (which is predefined already in $controller_node_internal, and the internal
# interface for compute nodes.
# In this example, eth0 is both the public and private interface for the
# controller.
# tunnel_ip allows you to create a network specifically for GRE tunneled
# traffic between compute and network nodes. Generally, you will want to
# use "ip" from the OpenStack node definitions above.

```

```

# This sets the IP for the private interface of compute and network nodes.

# Change build_server to the host name of your build node.
# Note that the hostname should be in all lowercase letters due to a
# limitation of Puppet (refer to http://projects.puppetlabs.com/issues/1168).
node build-server inherits build-node { }

# Change control-server to the host name of your control node. Note that the
# hostname should be in all lowercase letters due to a limitation of Puppet
# (refer to http://projects.puppetlabs.com/issues/1168).
node 'control-server' inherits os_base {
  class { 'control':
    tunnel_ip => '192.168.242.10'
  }
}

# Change compute-server01 to the host name of your first compute node.
# Note that the hostname should be in all lowercase letters due to a
# limitation of Puppet (refer to http://projects.puppetlabs.com/issues/1168).
node 'compute-server01' inherits os_base {
  class { 'compute':
    internal_ip => '192.168.242.21',
    tunnel_ip   => '192.168.242.21',
  }
}

### Repeat as needed ###
# Copy the compute-server01 line above and paste a copy here for each
# additional OpenStack node in your cluster. Be sure to replace the
# 'compute-server01' parameter with the correct host name for each
# additional node.

# Swift Proxy
# Adjust the value of swift_local_net_ip to match the IP address
# of your first Swift proxy node. It is generally not necessary
# to modify the value of keystone_host, as it will default to the
# address of your control node.
#node 'swift-proxy01' inherits os_base {
#  class {'openstack::swift::proxy':
#    swift_local_net_ip => $swift_proxy_address,
#    keystone_host      => $controller_node_address,
#    swift_user_password => $admin_password,
#    swift_admin_tenant  => 'admin',
#    swift_admin_user    => 'admin',
#    swift_hash_suffix   => $swift_hash,
#  }
#}

# Swift Storage
# Modify the swift_local_net_ip parameter to match the IP address of
# your first Swift storage node. Modify the storage_devices parameter
# to set the list of disk devices to be formatted with XFS and used
# for Swift storage.
#node 'swift-storage01' inherits os_base {
#  class {'openstack::swift::storage-node':
#    swift_zone => '1',
#    swift_local_net_ip => '192.168.242.180',
#    storage_type => 'disk',
#    storage_devices => ['sdb','sdc','sdd'],
#    swift_hash_suffix => $swift_hash,
#  }
#}

### Repeat as needed ###
# Copy the swift-storage01 node definition above and paste a copy here for
# each additional OpenStack swift storage node in your cluster. Modify
# the node name, swift_local_net_ip, and storage_devices parameters
# accordingly.

### End repeated nodes ###

#####

```

```

### All parameters below this point likely do not need to be changed ###
#####

### Advanced Users Configuration ###
# These four settings typically do not need to be changed.
# In the default deployment, the build node functions as the DNS and static
# DHCP server for the OpenStack nodes. These settings can be used if
# alternate configurations are needed.
$node_dns      = "${cobbler_node_ip}"
$ip            = "${cobbler_node_ip}"
$dns_service   = "dnsmasq"
$dhcp_service  = "dnsmasq"
$time_zone     = "UTC"

# Enable network interface bonding. This will only enable the bonding module
# in the OS, it won't actually bond any interfaces. Edit the networking
# interfaces template to set up interface bonds as required after setting
# this to true should bonding be required.
#$interface_bonding = 'true'

# Enable ipv6 router advertisement.
#$ipv6_ra = '1'

# Configure the maximum number of times mysql-server will allow
# a host to fail connecting before banning it.
$max_connect_errors = '10'

### Puppet Parameters ###
# These settings load other puppet components. They should not be changed.
import 'cobbler-node'
import 'core'

## Define the default node, to capture any un-defined nodes that register
## Simplifies debug when necessary.

node default {
  notify{"Default Node: Perhaps add a node definition to site.pp": }
}

```



Installing OpenStack

This chapter contains the following sections:

- [Summary of Steps for Installing OpenStack with Cisco OSI, page 21](#)
- [Cisco OSI Puppet Manifest Files, page 22](#)
- [Creating the Build Server, page 23](#)
- [Customizing the Build Server, page 26](#)
- [Building the Control and Compute Nodes, page 27](#)
- [Rerunning Puppet on the Build Node and Enabling Puppet to Run as an Agent, page 29](#)
- [Testing the OpenStack Deployment, page 29](#)

Summary of Steps for Installing OpenStack with Cisco OSI

- Step 1** Ensure that your environment meets the system requirements for Cisco OSI, as described in [System Requirements for Cisco OSI, on page 5](#).
- Step 2** Review the guidelines for installing OpenStack with Cisco OSI, as described in [Guidelines and Limitations of Cisco OSI, on page 2](#).
- Step 3** Create the build server that will serve as the build node for the OpenStack cluster.
You can perform this step in one of the following ways:
- By using the Cisco OSI install script, as described in [Creating the Build Server with the Cisco OSI Install Script, on page 23](#).
 - Manually by running each command in the script separately, as described in [Creating the Build Server Manually, on page 24](#).
- Step 4** Customize the build server, as described in [Customizing the Build Server, on page 26](#).
- Step 5** Build the control and compute nodes.
You can perform this step in one of the following ways:

- Build each control and compute node individually, as described in [Building the Control and Compute Nodes Individually](#), on page 27.
- Build the control and compute nodes with the Clean Node script, as described in [Building the Control and Compute Nodes with the clean_node.sh Script](#), on page 28. This script does the following:
 - Cleans out any existing certificate info from Puppet
 - Cleans out any existing cobbler node definitions for this same node name
 - Kicks off the booting of the nodes in Cisco UCS.
- Build the control and compute nodes with the Reset Nodes script, as described in [Building the Control and Compute Nodes with the reset_nodes.sh Script](#), on page 28. This script does the following:
 - Builds all the nodes defined in the `cobbler-node.pp` file.
 - Re-runs the Puppet apply and Puppet plugin download steps for the build node.

Step 6 Rerun Puppet on the build node to enable Puppet to run as an agent, as described in [Rerunning Puppet on the Build Node and Enabling Puppet to Run as an Agent](#), on page 29.

Step 7 Test the OpenStack deployment, as described in [Testing the OpenStack Deployment](#), on page 29.

Cisco OSI Puppet Manifest Files

The following table provides an overview of the files in the `/etc/Puppet/manifests` directory:

File Name	Description
<code>clean_node.sh</code>	Shell script that wraps several Cobbler and Puppet commands for ease of use when building and rebuilding the nodes of the OpenStack cluster.
<code>Cobbler-node.pp</code>	Manifest that manages the deployment of Cobbler, which supports the booting of additional servers into your environment.
<code>core.pp</code>	Manifest that defines the core definitions for the deployment of OpenStack services.
<code>modules.list</code>	File that provides a list of the Puppet modules.
<code>Puppet-modules.py</code>	Python script that downloads and installs all the necessary Puppet modules to deploy an OpenStack environment.
<code>reset_build_node.sh</code>	Shell script that deletes all configuration on your build node except the <code>site.pp</code> file and, therefore, effectively performs a complete clean out of your build node.

File Name	Description
reset_nodes.sh	Shell script that wraps around clean_node.sh to rebuild your entire cluster quickly with one command.
site.pp.example	<p>Manifest that contains the user modifiable components and defines the various parameters that must be set to configure the OpenStack cluster, including the Puppet Master and Cobbler setup on the build server.</p> <p>Note You must copy <code>site.pp.example</code> to <code>site.pp</code> and modify the settings for your environment.</p>

Creating the Build Server

Creating the Build Server with the Cisco OSI Install Script

The server that you use for the build node can be a physical server or a virtual machine (VM).

Before You Begin

- Ensure that the server you plan to use for the build node meets the minimum server requirements.
- Set your proxy configuration.

Step 1

Install Ubuntu 12.04 LTS on the server, as follows:

- Deploy no less than a minimal installation with **openssh-server**.
- Configure the network interface on the OpenStack cluster management segment with a static IP address.
- When partitioning the storage, choose a partitioning scheme that provides at least 15 GB of free space under `/var`, as the installation packages and ISO images that are used to deploy OpenStack are cached there.

Step 2

When the installation is complete, log in as root.

sudo -H bash

Step 3

If your environment includes a proxy server and you set your proxy information in your `.profile` or in a file such as `/etc/environment`, configure the following:

- `http_proxy`
- `https_proxy`
- A `no_proxy` command

The following is an example of the proxy configuration:

```
http_proxy=http://your-proxy.address.com:80/
https_proxy=https://your-https-proxy.address.com:443/
```

```
no_proxy=your-build-node-name,*yourbuild.domain.name,127.0.0.1,127.0.1.1,localhost
```

Step 4 If your environment does not include a proxy, copy and paste the following install script on the command line:

```
curl -s -k -B https://raw.githubusercontent.com/CiscoSystems/grizzly-manifests/multi-node/install_os_puppet |
/bin/bash
```

Step 5 If your environment includes a proxy, copy and paste the following install script on the command line:

```
https_proxy=http://proxy.example.com:80/ curl -s -k -B
https://raw.githubusercontent.com/CiscoSystems/grizzly-manifests/multi-node/install_os_puppet > install_os_puppet
chmod +x install_os_puppet
./install_os_puppet -p http://proxy.example.com:80/
```

What to Do Next

Customize the build server for your OpenStack environment.

Creating the Build Server Manually

The server that you use for the build node can be a physical server or a virtual machine (VM).

Before You Begin

- Ensure that the server you plan to use for the build node meets the minimum server requirements.
- Set your proxy configuration.

-
- Step 1** Install Ubuntu 12.04 LTS on the server, as follows:
- Deploy no less than a minimal installation with **openssh-server**.
 - Configure the network interface on the OpenStack cluster management segment with a static IP address.
 - When partitioning the storage, choose a partitioning scheme that provides at least 15 GB of free space under `/var`, as the installation packages and ISO images that are used to deploy OpenStack are cached there.
- Step 2** When the installation is complete, log in as root.
sudo -H bash
- Step 3** If your environment includes a proxy server and you set your proxy information in your `.profile` or in a file such as `/etc/environment`, configure the following:
- `http_proxy`
 - `https_proxy`
 - A `no_proxy` command

The following is an example of the proxy configuration:

```
http_proxy=http://your-proxy.address.com:80/
https_proxy=https://your-https-proxy.address.com:443/
no_proxy=your-build-node-name,*yourbuild.domain.name,127.0.0.1,127.0.1.1,localhost
```

- Step 4** Install pending security updates (if any) and the tools necessary for the installer to run, such as Puppet, git, and ipmitool.
- ```
apt-get update && apt-get dist-upgrade -y && apt-get install -y puppet git ipmitool
```

**Note** You might need to restart the server after you apply these updates.

- Step 5** Get the Cisco OSI example manifests from the [grizzly-manifests GitHub repository](#) branch that most closely matches your topology plans.  
If you need a version other than the default multi-node branch, you must include the **git checkout** tag whether your system has a proxy server or not.

The following example is for a system that does not use a proxy server and gets the example manifests from the multi-node branch.

```
git clone https://github.com/CiscoSystems/grizzly-manifests ~/cisco-grizzly-manifests/
cd ~/cisco-grizzly-manifests
git checkout -q g.0
```

The following example is for a system that uses a proxy server and gets the example manifests from the multi-node branch.

```
https_proxy=http://proxy.example.com:80 git clone https://github.com/CiscoSystems/grizzly-manifests
~/cisco-grizzly-manifests/
cd ~/cisco-grizzly-manifests
https_proxy=http://proxy.example.com:80
git checkout -q g.0
```

- Step 6** Copy the puppet manifests from `~/cisco-grizzly-manifests/manifests/` to `/etc/puppet/manifests/`.
- ```
cp ~/cisco-grizzly-manifests/manifests/* /etc/puppet/manifests
```

- Step 7** Copy the puppet templates from `~/cisco-grizzly-manifests/templates/` to `/etc/puppet/templates/`.
- ```
cp ~/cisco-grizzly-manifests/templates/* /etc/puppet/templates
```

- Step 8** Get the Cisco OSI puppet modules from the GitHub repository.  
The following example is for a system without a proxy server:

```
(cd /etc/puppet/manifests; python /etc/puppet/manifests/puppet-modules.py)
```

The following example is for a system with a proxy server:

```
(cd /etc/puppet/manifests; http_proxy=http://proxy.example.com:80
https_proxy=http://proxy.example.com:80
python /etc/puppet/manifests/puppet-modules.py)
```

## What to Do Next

Customize the build server for your OpenStack environment.

# Customizing the Build Server

**Step 1** Copy `site.pp.example` to `site.pp`.  
`cp /etc/puppet/manifests/site.pp.example /etc/puppet/manifests/site.pp`

**Step 2** Open `site.pp` and edit the settings and variables as appropriate for your environment.  
`vi /etc/puppet/manifests/site.pp`

**Step 3** Activate the manifest with **puppet apply**.  
`puppet apply -v /etc/puppet/manifests/site.pp`

When the **puppet apply** command runs, the Puppet client on the build server follows the instructions in the `site.pp` and `cobbler-node.pp` manifests, and configures the following programs on the build server:

- Network Time Protocol daemon (NTPD)—A time synchronization server used on all OpenStack cluster nodes to ensure that time throughout the cluster is correct.
- tftpd-hpa—A TFTP server that is part of the PXE boot process that occurs when OpenStack nodes boot up.
- dnsmasq—A DNS and DHCP server that is part of the PXE boot process that occurs when OpenStack nodes boot up.
- Cobbler—An installation and boot management daemon that manages the installation and booting of OpenStack nodes.
- apt-cacher-ng—A caching proxy for package installations that speeds up the package installation on the OpenStack nodes.
- Nagios—An infrastructure monitoring application that monitors the servers and processes of the OpenStack cluster.
- Collectd—A statistics collection application that gathers performance and other metrics from the components of the OpenStack cluster.
- Graphite and Carbon—A real-time graphing system that parses and displays metrics and statistics about OpenStack.
- Apache—A web server that hosts the sites needed to implement the following web services:
  - MySQL
  - Graphite
  - Nagios
  - Puppet Master passenger version, which is a Ruby runtime environment for Apache that provides a faster way to run the Puppet Master application web services

The initial Puppet configuration of the build server takes several minutes to complete as it downloads, installs, and configures the software needed for these applications.

**Step 4** After the initial Puppet configuration has been completed, stage the Puppet plugins so they can be accessed by the managed nodes.  
`puppet plugin download`

**Step 5** Verify that the nodes listed in the `site.pp` file were defined in Cobbler on the build server.

```
cobbler system list
 control-server
 compute-server01
 compute-server02
#
```

---

### What to Do Next

Use Cobbler to build the control and compute nodes.

## Building the Control and Compute Nodes

### Building the Control and Compute Nodes Individually

If you build the control and compute nodes individually, you can control the order in which the nodes are built. Therefore, if one or more nodes have a dependency on an application running on another node, you can ensure that the correct node is built first so that the other nodes do not fail. For example, you might want to build the control node first if you need Keystone fully configured on that node to ensure that other nodes can reach Keystone during their own installations.

With this procedure, you use the `clean_node.sh` script in Cobbler to build each of the nodes separately. The `clean_node.sh` script does the following:

- Configures Cobbler to PXE boot the specified node with the PXE options required to perform an automated install of Ubuntu.
- Uses Cobbler to power-cycle the node.
- Removes any existing client registrations for the node from Puppet, so that Puppet will treat this installation as a new installation.
- Removes any existing key entries for the node from the SSH known hosts database.
- Reboots the node.

### Before You Begin

You must complete the customization of the build server before you can build the control node.

---

**Step 1** Build your control node.

```
/etc/puppet/manifests/clean_node.sh {node_name}
```

Where `node_name` is the name of the control node in the `site.pp` file.

Once the installation finishes, the control node reboots. Then the script runs Puppet to pull and apply the control node configuration defined in the Puppet manifests on the build server. This step takes several minutes, as Puppet downloads, installs, and configures the various OpenStack components and support applications needed on the control node.

**Note** Setting up the control node might require more than one Puppet run, especially if there are proxies in the path, because some proxies can have issues with **apt-get** installations and updates. You can verify that the control node configuration has converged completely to the configuration defined in Puppet by looking at the log files in the `/var/log/syslog` directory on the control node.

**Step 2** (Optional) Observe the progress of the control node configuration as follows:

- View the automated install of Ubuntu on the KVM console of your control node.
- View the Puppet configuration run through the log files in the `/var/log/syslog` directory on the control node.

**Step 3** Run the following script for each compute node that you want to build:

```
/etc/puppet/manifests/clean_node.sh {node_name}
```

Where *node\_name* is the name of the compute node in the `site.pp` file.

As with the control node, building each compute node takes several minutes to complete.

## Building the Control and Compute Nodes with the `clean_node.sh` Script

The `clean_node.sh` script builds all the nodes defined in the `cobbler-node.pp` file. If you build the nodes with this script, it does the following:

- Cleans out any existing Puppet certificates.
- Cleans out any existing SSH parameters.
- Restarts the Cobbler build system.
- Uses Cobbler to power control the node according to the configuration in the `site.pp` file.

Run the `clean_node.sh` script.

```
for n in `cobbler system list`; do clean_node.sh $n ; done
```

## Building the Control and Compute Nodes with the `reset_nodes.sh` Script

The `reset_nodes.sh` script does the following:

- Builds all the nodes defined in the `cobbler-node.pp` file.
- Reruns the Puppet apply and Puppet plugin download steps for the build node.

Run the Reset Nodes script.

```
./reset_nodes.sh
```

---

## Rerunning Puppet on the Build Node and Enabling Puppet to Run as an Agent

This second Puppet run does the following:

- Removes all nodes from Cobbler.
- Reruns puppet to rebuild the Cobbler environment.
- Restarts all of the nodes in Cobbler.

### Before You Begin

Build all the control and compute nodes in your OpenStack cluster.

---

Run Puppet on the build node for a second time.

**puppet agent -t**

---

## Testing the OpenStack Deployment

After you build the nodes in the OpenStack cluster and the Puppet runs have completed on all nodes, you should test the OpenStack deployment.

### Step 1

Log into the OpenStack Horizon interface:

- a) In your browser, navigate to `http://control-node-IP/horizon/`
- b) Log into Horizon with the admin username and password in the `site.pp` file.  
If you did not change the defaults, the username is admin, and the password is Cisco123.

### Step 2

Load an image into the control node:

- a) Log into the console of the control node:
  - Username—localadmin
  - Password—ubuntu
- b) SU to root.
- c) In `/root/` run **source openrc**

d) Launch a test file in `/tmp/nova_test.sh`

---



## Using OpenStack

---

This chapter contains the following sections:

- [Deploying the First Virtual Machine Through a Script](#), page 31
- [Monitoring the Health of your System in Nagios](#), page 32
- [Viewing Statistics in Graphite](#), page 32
- [Additional Documentation](#), page 33

## Deploying the First Virtual Machine Through a Script

This procedure sets up a test Quantum network and launches the first VM on that network. The script that you use makes the following assumptions:

- That you will run the script as root, because the script creates keys and places them in the `/root/.ssh/` path. If you do not plan to run the script as root, you must change the path in the `create_vm` file for the `ssh-keygen` line.
- That you do not have existing files in the `/root/.ssh/` path. If you do have files in that path, the script prompts you to overwrite them. If you do not overwrite the existing keys, you might encounter a "permissions denied" error when you use SSH to access the test instance. If this error occurs, make sure that the key being referenced during the `nova keypair-add` step in the `create_vm` script is correct.

### Before You Begin

Complete clean Puppet runs on all OpenStack nodes.

---

#### Step 1

Clone the test repository.

**git clone https://github.com/CiscoSystems/quantum-l3-test**

After you complete this step, continue to Step 2 or follow the instructions in the `README.md` file.

#### Step 2

Navigate to the `quantum-l3-test` directory that you just cloned.

**cd quantum-l3-test**

**Step 3** Create the VM and set up the quantum network.

`./create_vm`

This script also runs **net\_setup** automatically.

**Step 4** When prompted by the script, do the following:

- Enter the network values for your public and private networks.
- Modify the default path for the Ubuntu Precise image to the path for your own local mirror.

**Step 5** Log into the instance.

To use the following example, replace *fixed-or-floating-ip-of-instance* with the appropriate value from your system.

```
ssh ubuntu@{fixed-or-floating-ip-of-instance}
```

**Step 6** (Optional) Reset the Quantum settings created by the script and relaunch the test VM.

This script method verifies your installation to ensure that it is working, sets up the appropriate Quantum networks, keys, uploads, and image, and then launches an instance.

```
./reset
./create_vm
```

---

## Monitoring the Health of your System in Nagios

The system and service health monitoring is included in the Puppet deployment through Cisco OSI. You can view status information about an OpenStack cluster with one or more VMs in Nagios.

### Before You Begin

Deploy at least one VM.

---

**Step 1** In your browser, navigate to the following URL: `http://ip-of-your-build-node/nagios3/`.

**Step 2** Log into Nagios:

- Username—admin
  - Password—Cisco123
- 

## Viewing Statistics in Graphite

You can view statistics about an OpenStack cluster with one or more VMs in Graphite.



### Before You Begin

Deploy at least one VM.

---

In your browser, navigate to the following URL: `http://ip-of-your-build-node:8190/`.

---

## Additional Documentation

The following documentation provides additional information about customizing, configuring, and maintaining your OpenStack cluster:

- [Cisco OpenStack documentation wiki](#)
- [OpenStack documentation](#) available on the OpenStack website

